Distributed Whiteboard Report

Alex Epstein &

October 15, 2022

1 Introduction

The system we present is a collaborative whiteboard where multiple users can simultaneously draw to a single consistent canvas. The network is structured as a single main server with multiple clients: the server maintains a canonical state of the whiteboard and performs user management functions.

From our testing we have established that the whiteboard performs properly under the required conditions. Multiple users are able to access and modify the whiteboard concurrently, and each user sees exactly the same drawings rendered on the whiteboard. Users are able to draw with the main shape tools (line, rectangle, circle, triangle), free line, and enter text. All drawings including text are colourable.

As additional features, we have implemented a user chat protocol, the ability to load and save a canvas to a file, and the ability for the admin user to kick a particular client by username. Similarly to the canvas itself, the user view of the chat and users list is apparently consistent across clients.

2 System Architecture & Communication Protocols

The whiteboard protocol was implemented using RMI. There are two main modules: client and server.

The client module is responsible for displaying the GUI (*ClientGUI*) and the rendered canvas (*InteractiveCanvas*), and hosts a remote object (*InteractiveCanvasManager*) responsible for communication with the server (both issuing and receiving method invocations).

The server is a single class (*RemoteWhiteboard*) which handles the canonical state of the whiteboard and user management functions. The canonical

whiteboard state is represented as a list of *Drawing* objects, each of which contain the necessary information to render them (origin, size, colour, etc.), along with a client signature and a timestamp.

An example drawing, involving a call from client to server, is shown in figure 1.
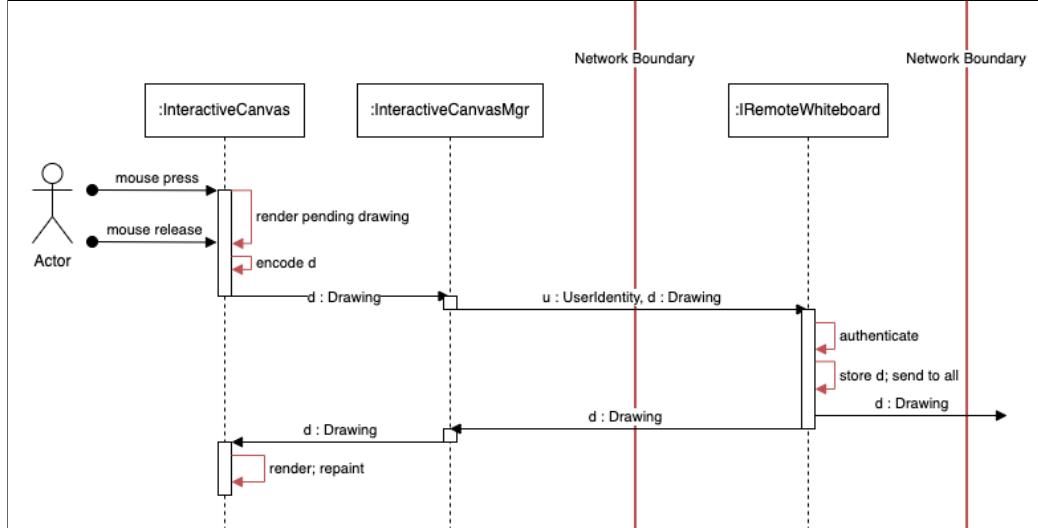

Figure 1: Example interaction diagram of an authenticated user submitting a drawing

In figure 1, a drawing of a rectangle is recognised by a mouse listener acting on the client's canvas. When the mouse is released, this triggers the canvas to send the drawing to its manager, which invokes the remote method in the server. In the server call, the canvas manager sends its *UserIdentity* along with the *Drawing*. The server subsequently verifies the user's identity, which consists of a username and password, then adds the drawing to its internal representation of the canvas. Then, the server issues a remote invocation to each of the client processes containing the drawing, which notifies them to render the new drawing.

At the client side, new drawings are received by the canvas manager, which adds them to a queue and notifies the canvas that drawings are pending. On receiving this notification, each drawing is drawn permanently to a flat image, which the client's canvas actually renders (the "canvas flat").

For user management operations, the procedure is similar, as shown in figure 2. Once a user has initiated a join request and been approved by the administrator, the server invokes methods in each client to notify them of a new user: similarly for when a user leaves, or a user sends a message in
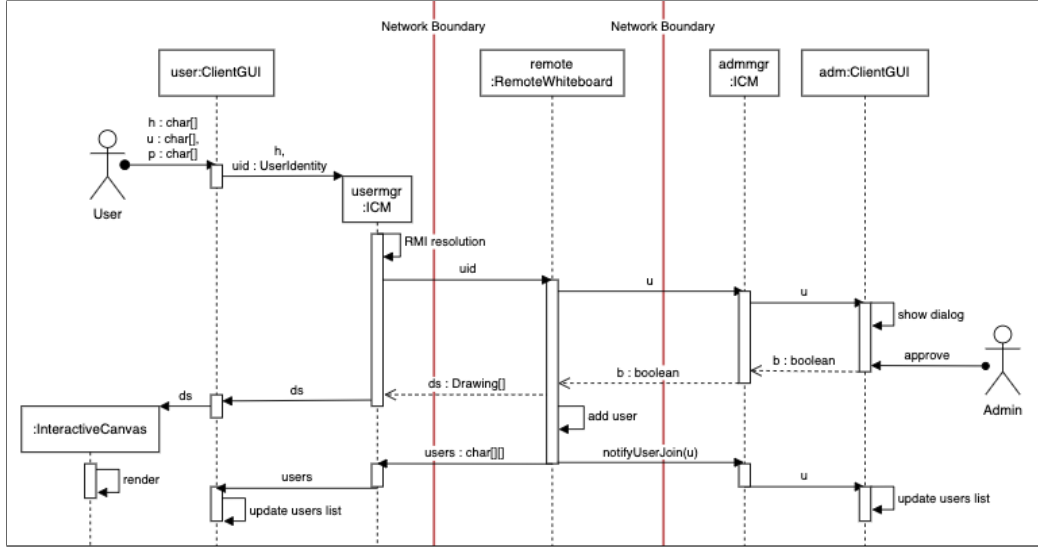
Figure 2: Example interaction diagram of a new user successfully joining the server

chat; the client uses one-way unicast and the server uses one-way multicast to update the global state.

3 Implementation Details

For the purposes of authentication, we have implemented a simple protocol which consists of the passing of *UserIdentity* objects between client and server. Each identity object contains the username and password of the sending client. Obviously, this is not secure, but the whiteboard is neither persistent nor a high-integrity system, and so plain-text was considered adequate. Users are also somewhat protected against spoofing by the fact that duplicate usernames are disallowed.

The *Drawing* package was designed to ensure small message sizes, and therefore throughput, and therefore the scalability of the system. Each drawing is represented by the minimum amount of information necessary to render it: in the case of simple shapes, a pair of integer coordinates is sufficient representation. In the case of free lines, a novel compression algorithm was written to minimise the amount of space required, reducing it by 80% in some cases while preserving all but the most granular details of the line.

Persistent storage of the canvas is accomplished locally on the client's machine, as a serialised array of drawings presented as a ".canvas" file. When saving, the client requests the drawings array from the server, and when

3

opening, the (admin) client deserialises the saved file and sends this to the server, which replaces the state with the saved state, changing the view for all clients. The performance bottleneck for this process is the number of *FreeLine* objects in the array, as the size of the other drawings is trivial. However, this has been somewhat mitigated, as shown above.

We now evaluate the non-functional requirements of the system, considering consistency, scalability, and performance.

Since requests to the server are served in the order they are received, there is little possibility that a user will see an inconsistent state unless the notification from the server to the user is lost or interrupted, which is considered a rare event: in this case, the user can disconnect and reconnect to remedy their local state. If a draw request from client to server fails, the drawing client will observe that the figure is not rendered on their local canvas, and so can retry. The order that each draw request arrives to the server is considered canonical, as we cannot rely on clock synchronisation between the clients. In short, barring rare events, the single-server architecture of the system maintains consistency.

Scalability of the system, though not a major concern, is difficult to assess given that Java RMI's implicit use of multithreading is not specified or standardised. Since all clients use strict unicast communication to and from the server, the server itself is the bottleneck: for each drawing sent by one of the n clients, the server must echo that drawing to all n users, which degrades performance linearly with large n - if each user sends m drawings per second, the server must make mn remote method calls per second, in addition to the increased overhead from user management.

However, this performance hit is mostly mitigated by the fact that *Drawing* objects are sent and stored at the server instead of raw images, greatly reducing message size. The server is *not responsible* for rendering each drawing as an image, which is a substantial consideration: the rendering process runs on the client machine, which experiences less load in general. Although this results in duplicated effort, minimising the server's required throughput is the critical factor, and so this results in greater health of the total system as it scales.